

[◀ Return to "Deep Learning" in the classroom](#)

Generate Faces

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Awesome submission, you have correctly implemented the DCGAN architecture and have used all the hyperparameters correctly. The results look great! Congratulations on passing this project and all the best 🍑

I suggest you read following articles to know how to improve results in GAN:

<https://github.com/soumith/ganhacks#how-to-train-a-gan-tips-and-tricks-to-make-gans-work>

<http://blog.otoro.net/2016/04/01/generating-large-images-from-latent-vectors/>

Required Files and Tests



The project submission contains the project notebook, called "dln_d_face_generation.ipynb".

Files correctly submitted.



All the unit tests in project have passed.

All the tests have passed.

Data Loading and Processing





The function `get_data_loader` should transform image data into resized, Tensor image types and return a `DataLoader` that batches all the training data into an appropriate size.

Dataloaders correctly used.



Pre-process the images by creating a `scale` function that scales images into a given pixel range. This function should be used later, in the training loop.

Well done implementing a general scale function with min and max values.

Build the Adversarial Networks



The Discriminator class is implemented correctly; it outputs one value that will determine whether an image is real or fake.

Excellent job at

- Using a sequence of convolutional layers using `conv2d` with strides to avoid making sparse gradients instead of max-pooling layers as they make the model unstable.
- Using `leaky_relu` and avoiding `ReLU` for the same reason of avoiding sparse gradients as `leaky_relu` allows gradients to flow backwards unimpeded.
- Using sigmoid as output layer.
- `BatchNorm` to avoid "internal covariate shift" as batch normalisation minimises the effect of weights and parameters in successive forward and back pass on the initial data normalisation done to make the data comparable across features.

Some suggestions for improvement:

- Use weight initialization: Xavier initialization is recommended as it helps model converge faster. A possible [implementation](#) in tensorflow is to pass `tf.contrib.layers.xavier_initializer()` as the value for the `kernel_initializer` parameter in `tf.layers.conv2d`



The Generator class is implemented correctly; it outputs an image of the same shape as the processed training data.

Same as the discriminator function: Excellent job at

- Implementing generator as "deconvolution" network with mostly the same key points as mentioned above.
- `Tanh` as the last layer of the generator output. This means that we'll have to normalise the input images to be between -1 and 1.



This function should initialize the weights of any convolutional or linear layer with weights taken from a normal distribution with a mean = 0 and standard deviation = 0.02.

Optimization Strategy



The loss functions take in the outputs from a discriminator and return the real or fake loss.

Excellent job implementing the loss of the model!



There are optimizers for updating the weights of the discriminator and generator. These optimizers should have appropriate hyperparameters.

Training and Results



Real training images should be scaled appropriately. The training loop should alternate between training the discriminator and generator networks.



There is not an exact answer here, but the models should be deep enough to recognize facial features and the optimizers should have parameters that help with model convergence.



The project generates realistic faces. It should be obvious that generated sample images look like faces.

The faces look great!



The question about model improvement is answered.

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)